

Agent-review outcomes under three evidence displays

12 gateway configurations × 12 frozen ExecRepoBench tasks; 48 planned episodes per condition

Generic review

executable / planned **44/48**

Correctness among executable outputs

38.6%

End-to-end success

35.4% 95% CI [27.1, 43.8]

Failure detection **54.5%**

Repair success **10.7%**

Uncertainty display

executable / planned **45/48**

Correctness among executable outputs

46.7%

End-to-end success

43.8% 95% CI [37.5, 50.0]

Failure detection **71.1%**

Repair success **27.6%**

Targeted guidance

executable / planned **45/48**

Correctness among executable outputs

35.6%

End-to-end success

33.3% 95% CI [25.0, 41.7]

Failure detection **68.2%**

Repair success **20.7%**

Uncertainty display – generic review

+8.3 pp 95% CI [−2.1, +18.8]

Targeted guidance – generic review

−2.1 pp 95% CI [−12.5, +8.3]